

Chapter 10

Bin Packing

Here we consider the classical BIN PACKING problem: We are given a set $I = \{1, \dots, n\}$ of *items*, where item $i \in I$ has *size* $s_i \in (0, 1]$ and a set $B = \{1, \dots, n\}$ of *bins* with *capacity* one. Find an assignment $a : I \rightarrow B$ such that the number of non-empty bins is minimal. As a shorthand, we write $s(J) = \sum_{j \in J} s_j$ for any $J \subseteq I$.

10.1 Hardness of Approximation

The BIN PACKING problem is NP-complete. More specifically:

Theorem 10.1. *It is NP-complete to decide if an instance of BIN PACKING admits a solution with two bins.*

Proof. We reduce from PARTITION, which we know is NP-complete. Recall that in the PARTITION problem, we are given n numbers $c_1, \dots, c_n \in \mathbb{N}$ and are asked to decide if there is a set $S \subseteq \{1, \dots, n\}$ such that $\sum_{i \in S} c_i = \sum_{i \notin S} c_i$. Given a PARTITION instance, we create an instance for BIN PACKING by setting $s_i = 2c_i / (\sum_{j=1}^n c_j) \in (0, 1]$ for $i = 1, \dots, n$. Obviously two bins suffice if and only if there is a $S \subseteq \{1, \dots, n\}$ such that $\sum_{i \in S} c_i = \sum_{i \notin S} c_i$. \square

This allows us to derive a lower bound on the approximability of BIN PACKING.

Corollary 10.2. *There is no ρ -approximation algorithm with $\rho < 3/2$ for BIN PACKING unless $P = NP$.*

10.2 Heuristics

We will show that there are constant factor approximations for BIN PACKING. Firstly we consider the probably most simple NEXT FIT algorithm, which can be shown to be 2-approximate. Secondly, we give the FIRST FIT DECREASING algorithm and show that it is 3/2-approximate. Thus, with the above hardness result, this is best-possible, unless $P = NP$.

Next Fit

The NEXT FIT algorithm works as follows: Initially all bins are empty and we start with bin $j = 1$ and item $i = 1$. If bin j has residual capacity for item i , assign item i to bin j , i.e., $a(i) = j$, and consider item $i + 1$. Otherwise consider bin $j + 1$ and item i . Repeat until item n is assigned.

Theorem 10.3. *NEXT FIT is a 2-approximation for BIN PACKING. The algorithm runs in $O(n)$ time.*

Proof. Let k be the number of non-empty bins in the assignment a found by NEXT FIT. Let k^* be the optimal number of bins. We show the slightly stronger statement that

$$k \leq 2 \cdot k^* - 1.$$

Firstly we observe the lower bound $k^* \geq \lceil s(I) \rceil$. Secondly, for bins $j = 1, \dots, \lfloor k/2 \rfloor$ we have

$$\sum_{i: a(i) \in \{2j-1, 2j\}} s_i > 1.$$

Adding these inequalities we get

$$\left\lfloor \frac{k}{2} \right\rfloor < s(I).$$

Since the left hand side is an integer we have that

$$\frac{k-1}{2} \leq \left\lfloor \frac{k}{2} \right\rfloor \leq \lceil s(I) \rceil - 1.$$

This proves $k \leq 2 \cdot \lceil s(I) \rceil - 1 \leq 2 \cdot k^* - 1$ and hence the claim. \square

The analysis is tight for the algorithm, which can be seen with the following instance with $2n$ items. For some $\varepsilon > 0$ let $s_{2i-1} = 2 \cdot \varepsilon$, $s_{2i} = 1 - \varepsilon$ for $i = 1, \dots, n$.

First Fit Decreasing

The algorithm NEXT FIT never considers bins again that have been left behind. Thus the wasted capacity therein leaves room for improvement. A natural way is FIRST FIT: Initially all bins are empty and we start with current number of bins $k = 0$ and item $i = 1$. Consider all bins $j = 1, \dots, k$ and place item i in the first bin that has sufficient residual capacity, i.e., $a(i) = j$. If there is no such bin increment k and repeat until item n is assigned. One can prove that FIRST FIT uses at most $k \leq \lceil 17/10 \cdot k^* \rceil$ many bins, where k^* is the optimal number.

There is a further natural heuristic improvement of FIRST FIT, called FIRST FIT DECREASING: Reorder the items such that $s_1 \geq \dots \geq s_n$ and apply FIRST FIT. The intuition behind considering large items first is the following: “Large” items do not fit into the same bin anyway, so we already use unavoidable bins and try to place “small” items into the residual space.

Theorem 10.4. *FIRST FIT DECREASING is a 3/2-approximation for BIN PACKING. The algorithm runs in $O(n^2)$ time.*

Proof. Let k be the number of non-empty bins of the assignment a found by FIRST FIT DECREASING and let k^* be the optimal number.

Consider bin number $j = \lceil 2/3k \rceil$. If it contains an item i with $s_i > 1/2$, then each bin $j' < j$ did not have space for item i . Thus j' was assigned an item i' with $i' < i$. As the items are considered in non-increasing order of size we have $s_{i'} \geq s_i > 1/2$. That is, there are at least j items of size larger than $1/2$. These items need to be placed in individual bins. This implies

$$k^* \geq j \geq \frac{2}{3}k.$$

Otherwise, bin j and any bin $j' > j$ does not contain an item with size larger than $1/2$. Hence the bins $j, j+1, \dots, k$ contain at least $2(k-j) + 1$ items, none of which fits into the bins $1, \dots, j-1$. Thus we have

$$\begin{aligned} s(I) &> \min\{j-1, 2(k-j) + 1\} \\ &\geq \min\{\lceil 2/3k \rceil - 1, 2(k - (2/3k + 2/3)) + 1\} \\ &= \lceil 2/3k \rceil - 1 \end{aligned}$$

and $k^* \geq s(I) > \lceil 2/3k \rceil - 1$. This even implies

$$k^* \geq \left\lceil \frac{2}{3}k \right\rceil \geq \frac{2}{3}k$$

and hence the claim. \square

10.3 Asymptotic Polynomial Time Approximation Scheme

With the hardness result that there is no approximation algorithm for BIN PACKING with guarantee better than $3/2$, unless $P = NP$, we do not have to search for a PTAS (or even an FPTAS). However, notice that the reduction used that the optimal number of bins is “small”, such as 2 or 3. It is plausible that, in “practical” instances, the optimal number k^* of bins grows as the number of items grows. Maybe we can do better for those instances.

This leads us to define: An *asymptotic polynomial time approximation scheme* (APTAS) is a family of algorithms, such that for any $\varepsilon > 0$ there is a number k' and a $(1 + \varepsilon)$ -approximation algorithm, whenever $k^* \geq k'$. For BIN PACKING such a family exists. However, the involved running times are rather high, even though polynomial in n .

Theorem 10.5. *For any $0 < \varepsilon \leq 1/2$ there is an algorithm that runs in time polynomial in n and finds an assignment having at most $k \leq (1 + \varepsilon) \cdot k^* + 1$ many bins.*

Lemma 10.6. *Let $\varepsilon > 0$ and $d \in \mathbb{N}$ be constants. For any instance of BIN PACKING where $s_i \geq \varepsilon$ and $|\{s_1, \dots, s_n\}| \leq d$, there is a polynomial time algorithm that solves it optimally.*

Proof. The number of items in a bin is bounded by $m := \lfloor 1/\varepsilon \rfloor$. Therefore, the number of different assignments for one bin is bounded by $r = \binom{m+d}{m}$, which is a (large) constant. There are at most n bins used and therefore, the number of feasible assignments is bounded by $p = \binom{n+r}{r}$. This is a polynomial in n . Thus we can enumerate all assignments and choose the best one to give an optimum solution. \square

Lemma 10.7. *Let $\varepsilon > 0$ be a constant. For any instance of BIN PACKING where $s_i \geq \varepsilon$, there is a $(1 + \varepsilon)$ -approximation algorithm.*

Proof. Let I be the given instance. Sort the n items by increasing size and partition them into $g = \lceil 1/\varepsilon^2 \rceil$ many groups each having at most $q = \lfloor n\varepsilon^2 \rfloor$ many items. Notice that two groups may contain items of the same size.

Construct an instance J by rounding up the size of each item to the size of the largest item in its group. Instance J has at most g many different item sizes. Therefore, we can find an optimal assignment for J by invoking Lemma 10.6. This is clearly a feasible assignment for the original item sizes.

Now we show that $k^*(J) \leq (1 + \varepsilon)k^*(I)$: We construct another instance J' by rounding down the size of each item to the smallest item size in its group. Clearly $k^*(J') \leq k^*(J)$.

The crucial observation is that an assignment for instance J' yields an assignment for all but the largest q items of the instance J . Therefore

$$k^*(J) \leq k^*(J') + q \leq k^*(I) + q.$$

To finalize the proof, since each item has size at least ε , we have $k^*(I) \geq n \cdot \varepsilon$ and $q = \lfloor n\varepsilon^2 \rfloor \leq \varepsilon \cdot k^*(I)$. Hence

$$k^*(J) \leq (1 + \varepsilon) \cdot k^*(I)$$

and the claim is established. \square

Proof of Theorem 10.5. Let I denote the given instance and I' the instance after discarding the items with size less than ε from I . We can invoke Lemma 10.7 and find an assignment which uses at most $k(I') \leq (1 + \varepsilon) \cdot k^*(I')$ many bins. By using FIRST FIT, we assign the items with sizes less than ε into the solution found for instance I' . We use additional bins if an item does not fit into any of the bins used so far.

If no additional bins are needed, then our assignment uses $k(I) \leq (1 + \varepsilon) \cdot k^*(I') \leq (1 + \varepsilon) \cdot k^*(I)$ many bins. Otherwise, all but the last bin have residual capacity less than ε . Thus $s(I) \geq (k(I) - 1)(1 - \varepsilon)$, which is a lower bound for $k^*(I)$. Thus we have

$$k(I) \leq \frac{k^*(I)}{1 - \varepsilon} + 1 \leq (1 + 2\varepsilon) \cdot k^*(I) + 1,$$

where we have used $0 < \varepsilon \leq 1/2$. \square