

Database

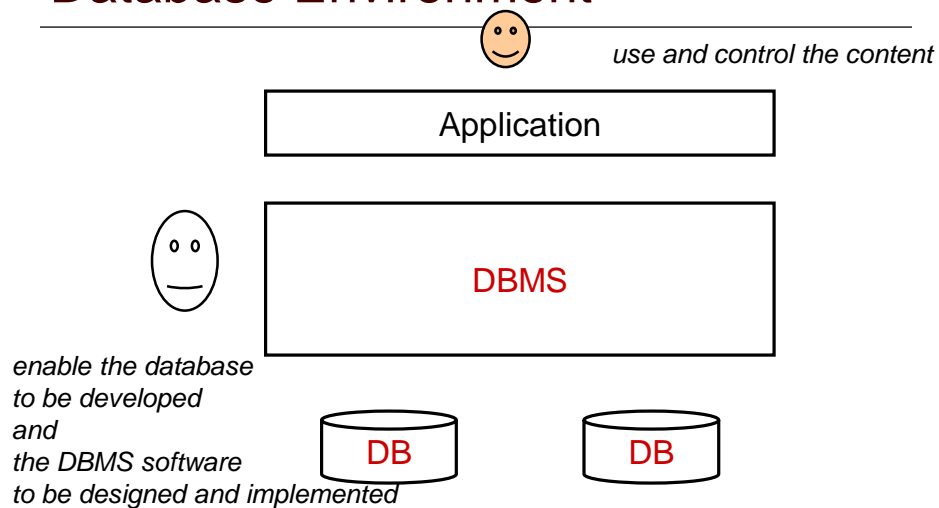
(some reminders)

Vu Tuyet Trinh

trinhvt@soict.hust.edu.vn

Department of Information Systems
School of Information Technology and Communication
Hanoi University of Science and Technology

Database Environment





Database

*A shared collection of related data
designed to meet the information needs
of an organisation*

- Logically coherent
- Internally consistent
- Specific purpose
- Representation of the real world

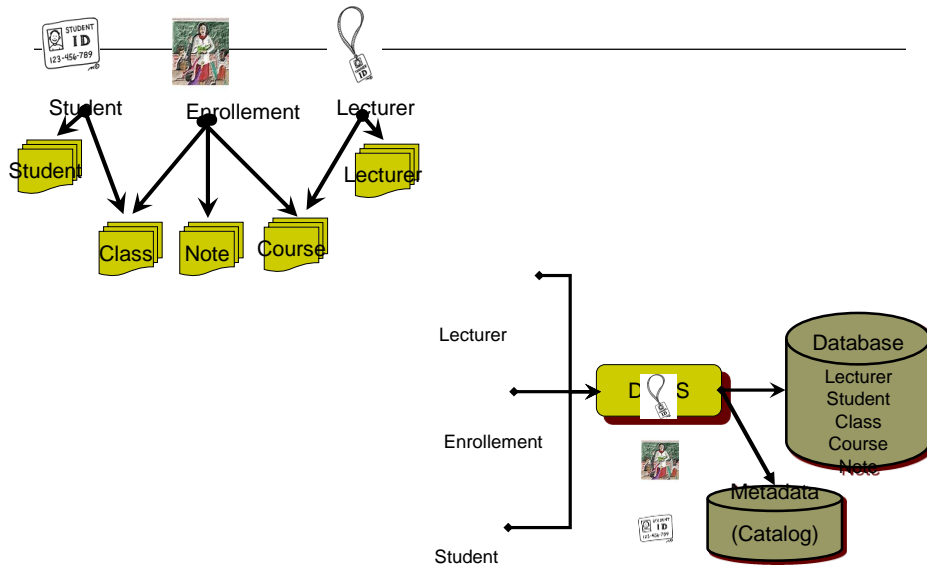


Database Management System

*A software to facilitate the creation and
maintenance of a database*

- Defining ~ specifying *types* of data
- Constructing ~ storing & populating
- Manipulating ~ querying, updating, reporting

File vs. Database



Database Advantages

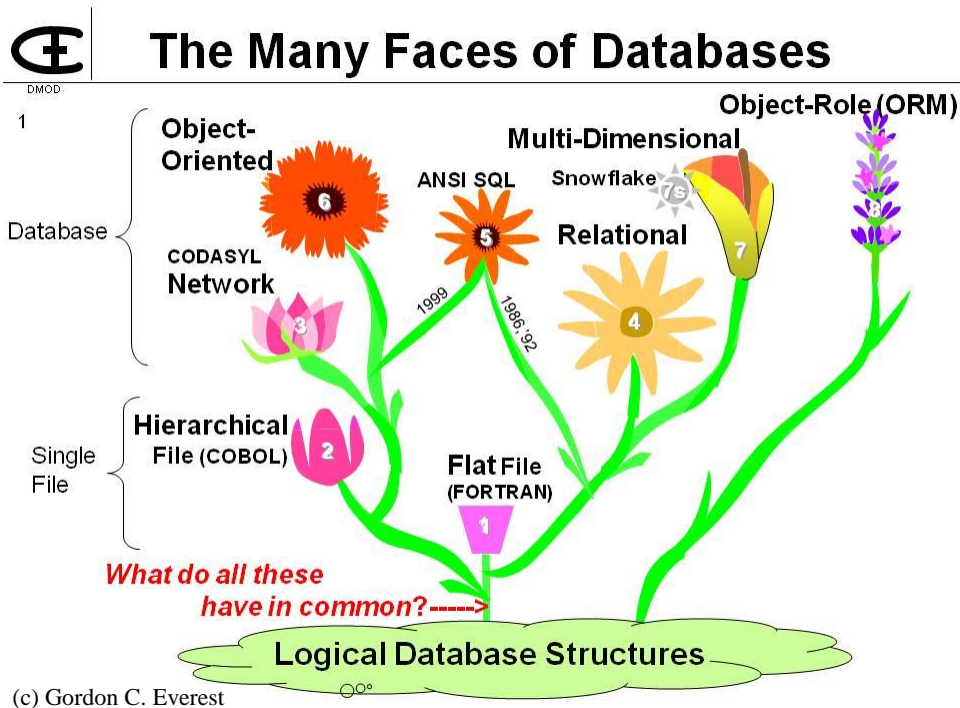
- ❑ Controlled redundancy
 - consistency of data & integrity constraints
- ❑ Integration of data
 - self-contained & represents semantics of application
- ❑ Data and operation sharing
 - multiple interfaces
- ❑ Flexibility
 - data independence
 - data accessibility
 - reduced program maintenance
- ❑ Services & Controls
 - security & privacy controls
 - backup & recovery
 - enforcement of standards
- ❑ Ease of application development

Data Models

- a set of concepts used to describe the **database structure**
 - data types
 - constraints
- Some existing database models
 - Hierarchical model
 - Network model
 - Relational model
 - Object-Oriented model

"More than **90%** of current **database applications** are built on **relational database systems** which uses **relational model** as its underlying data model"*

* R. Elmasri and S. Navathe. Fundamentals of Database Systems



Model vs. Schema vs. Instance

- Data Model
 - set of concepts used to describe the structure of a database:
data types, relationships, constraints, semantics, ...
 - tool for data abstraction
- Schema
 - data structure fulfilled all features of the parts of the real world
which is of interest to the users
- Instance
 - Data itself

Example

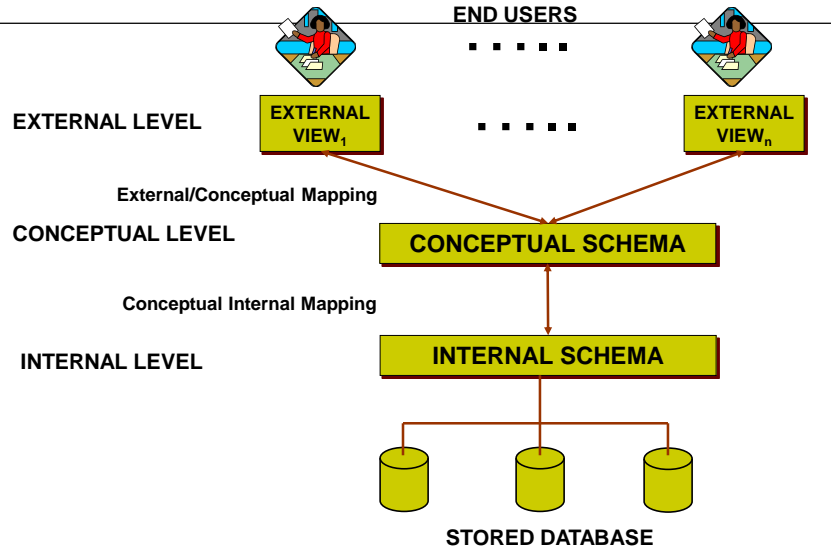
- Data Model

```
type <type_name> = record
    <field_name> : <data_type>;
    <field_name> : <data_type>;
    ...
end;
```
- Schema

```
type student = record
    ID : string;
    fullName: string;
    Birthday: date;
    Address: string ;
    Class: string;
end;
```
- Instance

(« Stud001 », « Nguyen », 1/4/1983, «1 Dai Co Viet », « 1F VN K50 »)

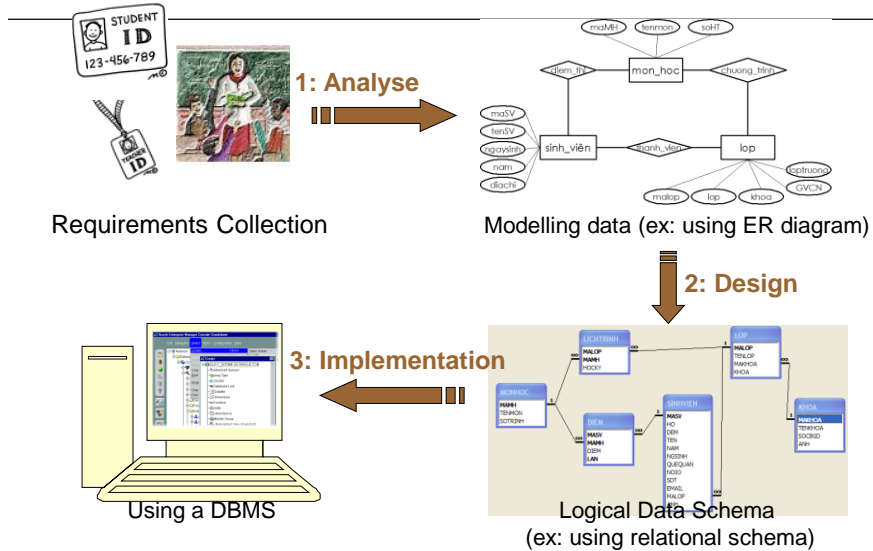
3-tier Schema Model (ANSI-SPARC Architecture)



Database Design

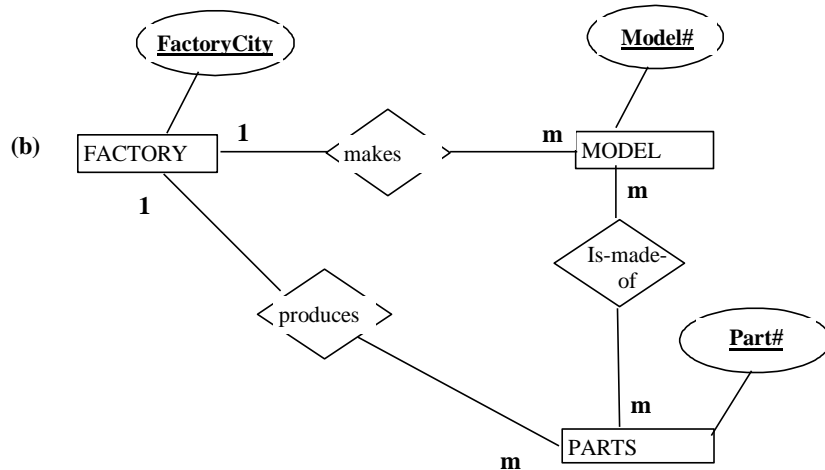
- Extended Entity Relationship
 - Top Down
 - Conceptual/Abstract View
- Functional Dependencies
 - Bottom Up
 - Implementation View
 - Synthesise relations

Process of Database Design



ER Model

- wellsuited to data modelling for use with databases
 - easy to represent and explain
 - readily translated to relations.
- Basic concepts
 - Attribute
 - represent a property/characteristic of an object in real world
 - Entity
 - defined as a set of attributes
 - Entity Set
 - Set of all entity instances of the same entity type
 - Relationship Set
 - Set of all relationship instances of the same relationship type



| <u>stud no</u> | name | tutor | roomno | <u>course no</u> | labmark | subject |
|----------------|--------|-------|--------|------------------|---------|----------|
| s1 | jones | bush | 2.26 | cs250 | 65 | prog |
| s1 | jones | bush | 2.26 | cs260 | 80 | graphics |
| s1 | jones | wibby | 2.26 | cs270 | 47 | elecs |
| s2 | brown | kahn | IT206 | cs250 | 67 | prog |
| s2 | brown | kahn | IT206 | cs270 | 65 | elecs |
| s3 | smith | goble | 2.82 | cs270 | 49 | comms |
| s4 | blogg | goble | 2.82 | cs280 | 50 | design |
| s5 | jones | zobel | 2.34 | cs250 | 0 | prog |
| s6 | peters | kahn | A17 | cs250 | 2 | prog |
| null | null | capon | A14 | null | null | null |
| null | null | null | null | cs290 | null | specs |

F

studno → name, tutor

tutor → roomno

roomno → tutor

courseno → subject

studno, courseno → labmark

F+

studno, courseno → name *partial*

studno → roomno *transitive*

Using functional dependencies to...Synthesise relations

$\text{studno} \rightarrow \text{studno}$
 $\text{studno} \rightarrow \text{familyname}$
 $\text{studno} \rightarrow \text{givenname}$
 $\text{studno} \rightarrow \text{hons}$
 $\text{studno} \rightarrow \text{tutor}$
 $\text{studno} \rightarrow \text{slot}$
 $\text{studno} \rightarrow \text{year}$

STUDENT

(studno, givenname, familyname, hons, tutor, slot, year)

$\text{studno, courseno} \rightarrow \text{labmark}$
 $\text{studno, courseno} \rightarrow \text{exammark}$

ENROL(studno, courseno, labmark, exammark)

$\text{courseno} \rightarrow \text{courseno}$
 $\text{courseno} \rightarrow \text{subject}$
 $\text{courseno} \rightarrow \text{equip}$

COURSE(courseno, subject, equip)

$\text{lecturer} \rightarrow \text{lecturer}$
 $\text{lecturer} \rightarrow \text{roomno}$
 $\text{lecturer} \rightarrow \text{appraiser}$
 $\text{roomno} \rightarrow \text{lecturer}$
 $\text{roomno} \rightarrow \text{roomno}$
 $\text{roomno} \rightarrow \text{appraiser}$

STAFF(lecturer, roomno, appraiser)

(L01, R01, 3)

(L02, R01, 4)

$\text{hons} \rightarrow \text{faculty}$
 $\text{hons} \rightarrow \text{hons}$

$\text{year} \rightarrow \text{year}$
 $\text{year} \rightarrow \text{yeartutor}$
 $\text{yeartutor} \rightarrow \text{year}$
 $\text{yeartutor} \rightarrow \text{yeartutor}$

YEAR(year, yeartutor)

SCHOOL(hons, faculty)

STUDENT_DETAILS

(studno, name, tutor, roomno, {courseno, labmark, subject})

$\text{studno} \rightarrow \text{name, tutor}$

$\text{courseno} \rightarrow \text{subject}$

$\text{tutor} \rightarrow \text{roomno, roomno} \rightarrow \text{tutor}$

$\text{studno, courseno} \rightarrow \text{labmark}$

STUDENT

(studno, name, tutor, roomno)

$\text{studno} \rightarrow \text{name, tutor}$

$\text{tutor} \rightarrow \text{roomno,}$

$\text{roomno} \rightarrow \text{tutor}$

ENROL (studno, courseno, subject, labmark)

$\text{courseno} \rightarrow \text{subject}$

$\text{studno, courseno} \rightarrow \text{labmark}$

COURSE (courseno, subject)

$\text{courseno} \rightarrow \text{subject}$

ENROL' (studno, courseno, labmark)

$\text{studno, courseno} \rightarrow \text{labmark}$

STUDENT (studno, name, tutor)

$\text{studno} \rightarrow \text{name, tutor}$

TUTOR (tutor, roomno)

$\text{tutor} \rightarrow \text{roomno}$

$\text{roomno} \rightarrow \text{tutor}$



Languages of DBMS

- Data Definition Language (DDL)
 - define the logical schema (relations, views, ...) and storage schema stored in a Data Dictionary
- Data Manipulation Language (DML)
 - Manipulative populate schema, update database
 - Retrieval querying content of a database
- Data Control Language (DCL)
 - permissions, access control, ...



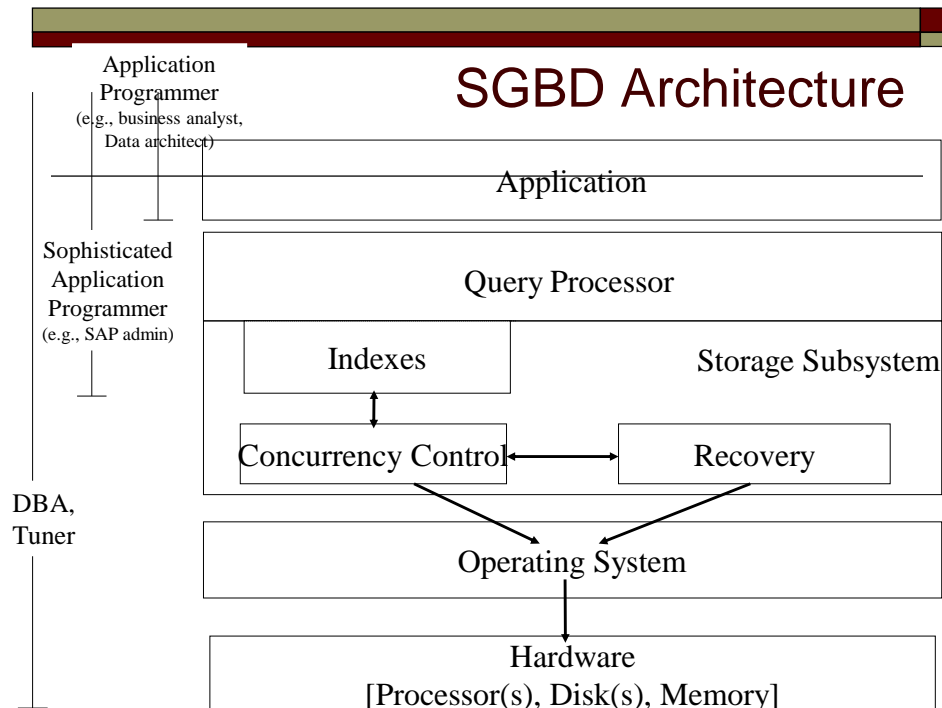
Data Manipulation Language

- Structured Query Language (SQL)
- A brief history
 - SQL 1
 - The first standard for SQL defined in 1986
 - adopted as an international by Standards Organisation (ISO) in 1987
 - SQL2
 - revised version of the processor (also called SQL 92).
 - adopted as the formal standard language for defining and manipulating relational database.
 - SQL 3
 - extension with additional features such as user-defined data types, triggers, user-defined functions and other Object Oriented features

SQL Retrieval Statement


```

SELECT [all|distinct]
        {*|{table.*|expr[alias]|view.*}
        [, {table.*|expr[alias]}]...}
FROM table [alias] [, table[alias]] ...
[WHERE condition]
[GROUP BY expr [,expr] ...]
[HAVING condition]
[ {UNION | UNION ALL | INTERSECT | MINUS }
    SELECT ... ]
[ORDER BY {expr|position} [ASC | DESC]
        [, expr|position} [ASC | DESC] .
    
```



Storage Management

- Responsible for storing and accessing data.
- Buffer manager
 - responsible for partitioning the available main memory into buffers
- File Management
 - responsible for interacting with file system



Indexing technique


- Search key
 - Any subset of the fields of a relation can be the search key
 - *Search key* may not be the *key* in relation
- Index
 - a collection of *k data entries*
 - supports efficient retrieval with a given key value *k*.

Classes of Indexes

- Primary vs. secondary: primary has primary key
- Clustered vs. unclustered: order of records and index approximately same
 - Alternative 1 implies clustered, but not vice-versa
 - A file can be clustered on at most one search key
- Dense vs. Sparse: dense has index entry per data value; sparse may “skip” some
 - Alternative 1 always leads to dense index
 - Every sparse index is clustered!
 - Sparse indexes are smaller; however, some useful optimizations are based on dense indexes

Query Processing

- Parser
 - verifying query syntax and semantic
- Optimizer
 - responsible for performing query plan transformation for the best evaluation
- Execution engine
 - Responsible for executing each of steps in the chosen query plan.
 - interacts with most of the other components of the DBMS



Relational Algebra Operations

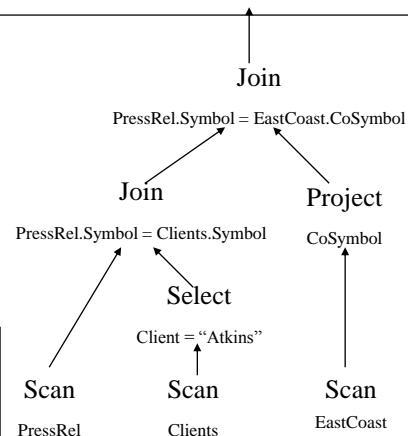
- Projection
- Selection
- Join
- Division
- Union
- Intersection
- Difference
- Cartesian product

27

Query Plans


- Data-flow graph of relational algebra operators
- *Typically:* determined by optimizer

```
SELECT *  
FROM PressRel p, Clients C  
WHERE p.Symbol = c.Symbol  
      AND c.Client = 'Atkins'  
      AND c.Symbol IN  
      (SELECT CoSymbol FROM EastCoast)
```



Transaction Management

- Ensure ACID properties
- The transaction processor performs the following tasks:
 - Logging
 - Recovery
 - Concurrency control




Transaction

- A sequence of read and write operations on data items that logically functions as one unit of work
 - Assuring data integrity and correction
- ACID Properties
 - **A**tomicity
 - **C**onsistency
 - **I**solation
 - **D**urability

} Concurrency Control
} Recovery

Transaction States



31

Transaction Management Interfaces

- Begin Trans
- Commit ()
- Abort()

- Savepoint Save()
- Rollback (savepoint)
(savepoint = 0 ==> Abort)

32

Concurrency Control

Objective:

- ensures that database transactions are performed concurrently without the concurrency violating the data integrity
- guarantees that no effect of committed transactions is lost, and no effect of aborted (rolled back) transactions remains in the related database.

Example

| | |
|---|---|
| <p>T0: read(A); A := A - 50; write(A); read(B); B := B + 50; write(B);</p> | <p>T1: read(A); temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; write(B);</p> |
|---|---|

33

Scheduling

| T0 | T1 | T0 | T1 | T0 | T1 |
|--|---|---|---|--|---|
| read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) | read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) | read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) |

(1)

(2)

(3)

34

Serializability

- A **schedule** of a set of transactions is a linear ordering of their actions
 - e.g. for the simultaneous deposits example:
R1(X) R2(X) W1(X) W2(X)
- A **serial schedule** is one in which all the steps of each transaction occur consecutively
- A **serializable schedule** is one which is equivalent to some serial schedule

Lock

- **Definition**
 - a synchronization mechanism for enforcing limits on access to DB in concurrent way.
 - one way of enforcing concurrency control policies
- **Lock types**
 - **Shared lock** (LS) readable but can not write
 - **Exclusive lock** (LX): read and write
 - UN(D): unlock
- **Compatibility**


| | LS | LX |
|----|-------|-------|
| LS | true | false |
| LX | false | false |

How can constraints be violated?

- ❑ Transaction bug
- ❑ DBMS bug
- ❑ Hardware failure
 - e.g., disk crash
- ❑ Data sharing
 - e.g., T1 and T2 in parallel

Failures

Events — Desired
 \ Undesired — Expected
 \ Unexpected





Recovery

- Maintaining the consistency of DB by ROLLBACK to the last consistency state.
- Ensuring 2 properties
 - Atomic
 - Durability
- Using LOG




Transaction Log

- A sequence of log record keeping trace of actions executed by DBMS
 - <start T>
 - Log the beginning of the transaction execution
 - <commit T>
 - transaction is already finished
 - <abort T>
 - Transaction is calcel
 - <T, X, v, w>
 - Transaction makes an update actio, before update X=v, after update x = w

-
- Read(A)
 - If A > 50 then display("so du hop le")
 - Else {
 A:=A+50
 =====→CRASH
 display ("ghi no tai khoan A")
}

Transaction Log

- Handled in main memory and put to external memory (disk) when possible





Checkpoint

- Definition:
 - moment where intermediate results and a log record are saved to disk.
 - being initiated at specified intervals
- Objective
 - minimize the amount of time and effort wasted when restart
 - the process can be restarted from the latest checkpoint rather than from the beginning.
- Log record
 - <checkpoint> or <ckpt>



Discussion

- Undo Logging
 - need to write to disk as soon transaction finishes
 - Access disk
- Redo Logging
 - need to keep all modified blocks in memory until commit
 - Use memory

Undo/Redo Logging Recovery Rules

- ❑ Backwards pass (end of log \rightarrow latest valid checkpoint start)
- ❑ Constructing set S of committed transactions
- ❑ undo actions of transactions not in S
- ❑ undo pending transactions
- ❑ follow undo chains for transactions in (checkpoint active list) – S
- ❑ Forward pass (latest checkpoint start \rightarrow end of log)
- ❑ redo actions of S transactions

